

Verifact

Check what a YouTube video claims, while you watch it

team admini

HackYeah Defence 2026



The correction arrives too late

THE GAP

Someone watches a twenty-minute video tonight.

- If one claim in it is false, the correction—if it ever comes— lands about two weeks from now.
- Nine times out of ten it never lands at all.
- By then they have already repeated it at work.

We move the check to the moment of exposure.

Community Notes: 14-day median, 90%+ never published · Harvard Misinformation Review, 2025 (external).

Why this is Defence

RELATION TO CATEGORY

Disinformation is not a content problem. It is a defence problem, because the target is a population rather than a system.

- Defence means protecting a population's picture of reality.
- The attack surface is the moment a claim is believed.
- The tool has to sit where the belief forms: the player.

The solution: Verifact

WHAT IT IS

A browser extension that checks a video's claims while you watch, and shows the evidence behind every verdict.

- Reads the video's captions, extracts the checkable claims, and drops opinions and predictions.
- Retrieves evidence per claim from reference, academic and fact-checking sources, in parallel.
- Adjudicates, then argues against itself before publishing a verdict.
- Fetches every cited page and confirms the quoted line is on it.

False and doubtful claims are marked on the timeline; the side panel carries the verdicts and their sources.

Demo

THE CHECK, LIVE ON A REAL VIDEO

Reading the captions. Finding the claims. Going to look.

- Markers appear on the scrubber as each claim is decided.
- Scrub to a red marker: the notice, then the source behind it.
- Every link opens a page that contains the quoted line.

Roughly half the practical score lives in these three minutes.

We go and look

INNOVATION · CITATION VERIFICATION

Models fabricate citations. We fetch every cited page and look.

False

02:13

The video states a claim the evidence contradicts.

EVIDENCE

National Archives

reference

“the quoted span, string-matched on the page”

A quote not on its cited page downgrades the verdict-by string matching, not a second model.

How it is argued, not just whether it is true

INNOVATION · RHETORICAL PASS

A claim can be accurate and still be deployed to mislead. That reading runs beside fact-checking, never inside it.

- It names the technique—emotional pressure, loaded language, faulty reasoning, selective evidence, conspiracy framing.
- Every signal must quote a verbatim transcript line or it is dropped; the prompt names the technique, never the position.

An empty list is the normal, correct answer.

We are not first

PRIOR ART

PopUpFactCheck, Factly, DeepVerify and Live Fact Checker already do in-video fact-checking. We get in front of that rather than being corrected.

- **What they do:** check claims while the video plays.
- **What none of them do:** verify a cited page contains the quoted line.
- **What none publish:** an accuracy figure.

The narrow claim is the defensible one. Never claim first, only, or that it “solves misinformation”.

Who pays

PRACTICAL APPLICABILITY

Checks out

individuals

Freemium. The viewer-side tool has to be free to spread.

Checks out

newsrooms and schools

Per-seat, where a check is traceable and part of a workflow.

Doubtful

institutions

DSA and AI Act reporting duties—the institutional wedge.

What runs today, what we would fix next

SHIPPED VS NOT

RUNS TODAY

- Browser extension and backend
- Streaming claim-level verdicts
- Citation verification by string match
- Rhetorical pass, CLI, labelled evals

NEXT

- Deploy beyond a local backend
- More languages, more than YouTube
- A local model to drop the per-check cost
- Score the whole video, not only claims

Everything on the left you just watched run. The right is what we would fix next.