

Verifact

Check what a YouTube video claims, while you watch it

team admini

HackYeah Defence 2026



The correction arrives too late

THE GAP

Someone watches a twenty-minute video tonight.

- If one claim in it is false, the correction—if it ever comes— lands about two weeks from now.
- Nine times out of ten it never lands at all.
- By then they have already repeated it at work.

We move the check to the moment of exposure.

Why this is Defence

RELATION TO CATEGORY

Disinformation is not a content problem. It is a defence problem, because the target is a population rather than a system.

- Defence means protecting a population's picture of reality.
- The attack surface is the moment a claim is believed.
- The tool has to sit where the belief forms: the player.

The solution: Verifact

WHAT IT IS

A browser extension that checks a video's claims while you watch, and shows the evidence behind every verdict.

- Reads the video's captions, extracts the checkable claims, and drops opinions and predictions.
- Retrieves evidence per claim from reference, academic and fact-checking sources, in parallel.
- Adjudicates, then argues against itself before publishing a verdict.
- Fetches every cited page and confirms the quoted line is on it.

False and doubtful claims are marked on the timeline; the side panel carries the verdicts and their sources.

Demo

THE CHECK, LIVE ON A REAL VIDEO

Reading the captions. Finding the claims. Going to look.

- Markers appear on the scrubber as each claim is decided.
- Scrub to a red marker: the notice, then the source behind it.
- Every link opens a page that contains the quoted line.

Roughly half the practical score lives in these three minutes.

We go and look

INNOVATION · CITATION VERIFICATION

Models fabricate citations. We fetch every cited page and look.

False

02:13

The video states a claim the evidence contradicts.

EVIDENCE

National Archives

reference

"the quoted span, string-matched on the page"

How it is argued, not just whether it is true

INNOVATION · RHETORICAL PASS

A claim can be accurate and still be deployed to mislead. That reading runs beside fact-checking, never inside it.

- It names the technique—emotional pressure, loaded language, faulty reasoning, selective evidence, conspiracy framing.
- Every signal must quote a verbatim transcript line or it is dropped; the prompt names the technique, never the position.

We are not first

PRIOR ART

PopUpFactCheck, Factly, DeepVerify and Live Fact Checker already do in-video fact-checking. We get in front of that rather than being corrected.

- **What they do:** check claims while the video plays.
- **What none of them do:** verify a cited page contains the quoted line.
- **What none publish:** an accuracy figure.

Who pays

PRACTICAL APPLICABILITY

Checks out

individuals

Freemium. The viewer-side tool has to be free to spread.

Checks out

newsrooms and schools

Per-seat, where a check is traceable and part of a workflow.

Doubtful

institutions

DSA and AI Act reporting duties—the institutional wedge.

What runs today, what we would fix next

SHIPPED VS NOT

RUNS TODAY

- Browser extension and backend
- Streaming claim-level verdicts
- Citation verification by string match
- Rhetorical pass, CLI, labelled evals

NEXT

- Deploy beyond a local backend
- More languages, more than YouTube
- A local model to drop the per-check cost
- Score the whole video, not only claims